RICE UNIVERSITY · DEPARTMENT OF ECONOMICS

AI in ECON 209: Policy, Tools, and Privacy

Applied Econometrics (ECON 209 004/005) · Fall 2026 · Prof. Jeff Dominitz
Prepared by Jessica Rood, Teaching Assistant | jr120@rice.edu
Companion to *Getting Set Up*. Read it when you have a moment — nothing here needs doing before the first lab.

The syllabus permits AI tools for coding on the problem sets, which raises three practical questions: where exactly the permitted line falls, which tools actually do what they claim, and what these companies do with what you type. This handout takes them in that order.

1 The course AI policy

The syllabus addresses this directly, and it's worth reading in full rather than in summary:

> "*Use of AI is permitted to assist you with your coding for the Problem Sets. The use of AI will not be allowed on the Problem Sets for interpreting the results of your data analysis, nor will use of AI be allowed for answering Problem Set questions that do not involve coding. Use of AI will not be allowed in any form for the exams.*"

The distinction is between **writing code** and **interpreting what the code produced**. And the no-AI rule covers **all three exams**, both midterms and the final.

The syllabus states plainly that **unauthorized use of AI is a Rice Honor Code violation**, and it requires the following signed statement with each problem set and exam:

> "*On my honor, I have neither given nor received any unauthorized aid on this [problem set / exam].*"

We'll talk through how the coding/interpreting line applies in practice during our first lab, with examples on both sides. If you're ever unsure whether a particular use is appropriate, please ask Prof. Dominitz or me *before* you submit rather than after. Where this handout and the syllabus differ, the syllabus governs; this is a TA-prepared guide to help you get started, not a change to course policy.

2 A note on agentic tools

"Agentic" describes a category of tool rather than a specific product. The difference from an ordinary chat window is that these tools work inside your project, reading your files, running your code, seeing what went wrong, and trying again, with you reviewing the changes. **Claude Code** (Anthropic), **Codex** (OpenAI), and **GitHub Copilot** all work this way and all integrate with VSCode.

One thing worth knowing before you install anything: for these products, the *free* tier is generally the chat version, not the agentic one. Claude Code is the clearest example — its VSCode extension requires a paid Claude plan (Pro or Max) or a pay-per-use Console account, so a free Claude account will let you use the website but not the tool inside your editor. Terms change often enough that each product's own pricing page is the authority; check it before assuming a free tier covers the in-editor version.

No paid tool is required for this course. Everything we do in lab runs without one, and nothing in your grade depends on which tool you use, or on using one at all.

GEMINI IS FREE TO YOU THROUGH RICE

Rice provides **Gemini** at no cost through your Rice Google account, and it will handle most of what you need in this course: explaining code, working through an error message, or answering "how do I do this in pandas." If you'd rather not sign up for anything new, this is a perfectly good place to start.

One limitation to know going in: Rice's license covers the **chat interface only**, so you won't be able to use Gemini inside VSCode. You'll be copying between a browser tab and your editor rather than having it work directly with your files. That's a difference in convenience rather than in what it can help you figure out.

WHAT I'LL BE USING

The ideas here apply to whichever tool you choose, but so you know what you'll see on screen: when I demonstrate something in lab or office hours, **I'll be using a paid version of Claude.** Everything I show can be done without one. If you're working in Codex, Copilot, or Gemini, the concepts should transfer even where the exact keystrokes don't. Do mention it during lab and I'll try to point out where your setup differs from what I'm showing.

Whichever you use, one habit is worth carrying through the semester: **be able to explain any number you report.** Where the estimate came from, what it assumes, and why you find it credible. That's a useful standard regardless of how the code got written, and it's the part of this work that stays with you.

KEEP AN AI JOURNAL

Keep a short running log of how you use AI in this course — what you asked it to do, where it saved you time, where it steered you wrong. A plain text file in your course folder is enough. Two reasons to bother: it keeps you honest about which parts of the work are yours, exactly the line the policy asks you to hold, and by December you'll have your own record of what these tools are good for, which beats anyone else's general advice because it's about how *you* work.

Using an agentic tool? Hand it the record-keeping. Ask it to maintain the log as you go — the decisions you made, where you changed direction, what went wrong and how you fixed it. The tool writes down *what happened*; you write the reflection at the end. The judgment stays yours, and you're not reconstructing a session from memory a month later.

3 Privacy: what these tools do with what you type

On personal accounts, the tools above use your conversations to train future models *by default*. You can turn that off, and it's worth doing: it takes about a minute per tool, and the setting isn't shown to you during signup. Your Rice account is the exception.

Tool	Default on a personal account	Where to turn it off
Gemini via your Rice account	**Not used for training**, and not reviewed by humans. This comes with Workspace for Education	Nothing to change. Just be sure you're signed in with your Rice account, not a personal one
Claude Free, Pro, Max	Chats and coding sessions used for training unless you opt out	Profile icon → **Settings** → **Privacy** → turn off "*Help improve Claude*"
ChatGPT Free, Plus, Pro	On by default on personal workspaces	**Settings** → **Data Controls** → turn off "*Improve the model for everyone.*" Temporary Chats are excluded either way
GitHub Copilot Free, Pro, Pro+	On by default since April 2026, including code from private repositories	`github.com/settings/copilot/features` → under **Privacy**, turn off data use for AI model training

Two things to know about what opting out does and doesn't do. It applies going forward, so it doesn't remove anything from models already trained. And it is **not the same as privacy**: providers still retain conversations for a period for safety and abuse monitoring. Part of why the Rice Gemini account is a reasonable default is that its protection is built in rather than something you have to go find and switch on.

> **IN PRACTICE, FOR THIS COURSE**
>
> Our datasets are public, so nothing you do in lab is sensitive. This matters more for what comes after: a thesis, an RA position with restricted data, an internship with a company's numbers. As a rule, don't paste restricted or confidential data into these tools; if we work with anything of that kind, I'll flag it.
>
> Menus move around. If a path above doesn't match what you see, look for **Privacy** or **Data Controls** in settings.

4 If you want to read further

Entirely optional, and nothing in the course depends on it. Scott Cunningham, an economist at Baylor and the author of *Causal Inference: The Mixtape*, has been writing about how these tools are actually being used in economics since December 2025 at causalinf.substack.com. Two posts are a useful orientation.

> **What this is · "Claude Code Changed How I Work (Part 1)"**
> *December 13, 2025* → *causalinf.substack.com/p/claude-code-changed-how-i-work-part*
>
> Cunningham distinguishes **agentic coding**, where a tool works inside your project and runs your code, from simply accepting whatever a chat window hands back — the distinction section 2 above is about. This is the first of a long series he is still adding to. *(Behind his paywall; you may only see a preview.)*

> **Where the limits are · "Which research mistakes matter for human capital?"**
> *July 27, 2026* → *causalinf.substack.com/p/which-research-mistakes-matter-for*
>
> He describes a coding error in an early paper that changed his results, and what the episode taught him about organizing files and code. He also leaves open a question worth considering: if these tools catch every mistake, is anything lost?

Questions? Email me at `jr120@rice.edu`, or raise it in lab — policy questions in particular are worth asking out loud, since other people usually have the same one. · ECON 209 004/005 · Fall 2026 · Prepared by Jessica Rood, TA; not an official University publication, and the syllabus governs where the two differ.